

Ismail Elezi

London, UK | ismail.elezi@gmail.com | scholar | LinkedIn | Twitter | Github | Web |

I am a Principal Research Scientist of Computer Vision, leading the multi-modality learning team in Huawei Noah's Ark lab in London. I worked before in deep metric learning, self and semi-supervised learning, open-world and long-tail detection, active learning, and generative models (GANs and diffusion models). Currently I am focused on researching new models in multi-modality learning (visual LLMs). I frequently publish in top tier vision (CVPR, ICCV, ECCV) and ML conferences (NeurIPS, ICML, ICLR). ***h-index: 15, citations: 1154.***

WORKING EXPERIENCE

Deep Learning Principal Researcher: Huawei <i>Leading a team working in multi-modality learning, VLMs and diffusion models.</i>	December 2024 - <i>London, UK</i>
Deep Learning Senior Researcher: Huawei <i>Build and leading a team in multi-modality learning and VLMs.</i>	April 2023 - November 2024 <i>London, UK</i>
Deep Learning Research Visitor: Argo AI <i>Worked on semi-supervised LiDAR data for object segmentation.</i>	June 2022 - September 2022 <i>Munich, Germany</i>
Alexander von Humboldt Postdoctoral Researcher: TUM <i>Research, mentored Ph.D. and master students, co-taught several courses.</i>	Oct 2020 - March 2023 <i>Munich, Germany</i>
Deep Learning Research Intern: NVIDIA <i>Resulted in papers accepted to CVPR and ICCV.</i>	Feb. 2020 - Sept. 2020 <i>Santa Clara, US</i>

EDUCATION

Ph.D. in Deep Learning <i>Ca' Foscari University of Venice, Venice, Italy</i>	Sep. 2016 – July 2020 <i>Grade: Distinction</i>
---	--

- Completed Ph.D. under the supervision of professors Marcello Pelillo and Thilo Stadelmann. Spent a year at TUM.

SELECTED PUBLICATIONS

- Alexandridis, **Elezi**, Deng, Nguyen, and Luo. Fractal calibration for long-tailed object detection. *CVPR25*
- Kaul, Ma, **Elezi**, and Deng. From attention to activation: Unravelling the enigmas of large language models. *ICLR25*
- Miles, Reddy, **Elezi**, and Deng. Velora: Memory efficient training using rank-1 sub-token projections. *NeurIPS24*
- Reddy*, **Elezi***, and Deng. G3dr: Generative 3d reconstruction in imagenet. *CVPR24*
- Miles, **Elezi**, and Deng. Vkd: Improving knowledge distillation using orthogonal projections. *CVPR24*
- Ma, **Elezi**, Deng, Dong, and Xu. Three heads are better than one: Complementary experts for long-tailed semi-supervised learning. *AAAI24*
- Seidenschwarz, Brasó, Serrano, **Elezi**, and Leal-Taixé. Simple cues lead to a strong multi-object tracker. *CVPR23*
- Elezi***, Seidenschwarz*, Wagner*, Vascon, Torcinovich, Pelillo, and Leal-Taixé. The group loss++: A deeper look into group loss for deep metric learning. *tPAMI23*
- Kocsis, Sukenik, Brasó, Niessner, Leal-Taixé, and **Elezi**. The unreasonable effectiveness of fully-connected layers for low-data regimes. *NeurIPS22*
- Fomenko, **Elezi**, Ramanan, Osep, and Leal-Taixé. Learning to discover and detect objects. *NeurIPS22*
- Elezi**, Yu, Anandkumar, Leal-Taixé, and Alvarez. Not all labels are equal: Rationalizing the labeling costs for training object detection. *CVPR22*

12. Choi, **Elezi**, Lee, Farabet, and Alvarez. Active learning for deep object detection via probabilistic modeling. *ICCV21*
13. Seidenschwarz, **Elezi**, and Leal-Taixé. Learning intra-batch connections for deep metric learning. *ICML21*
14. **Elezi**, Vascon, Torcinovich, Pelillo, and Leal-Taixé. The group loss for deep metric learning. *ECCV20*
15. Maximov*, **Elezi***, and Leal-Taixé. CIAGAN: conditional identity anonymization generative adversarial networks. *CVPR20*

* = equal contribution. For a full list of papers, please check my Google Scholar.

SKILLS

Programming: Python, PyTorch, Tensorflow, openCV, sklearn, Java, Matlab/Octave, C, C#, C++.

Languages: Albanian (Mother tongue), English (fluent), Italian (intermediate), German (beginner).

REVIEWING DUTIES

Area Chair: WACV 2021, NeurIPS 2025

Conferences: CVPR 2020, 2021*, 2022, 2023, 2024, 2025; ICCV 2021*, 2025; ECCV 2022, 2024; NeurIPS 2021; ICML 2022; ICLR 2024, 2025; IJCAI 2021; BMVC 2019, 2020; ACCV 2020*; WACV 2022. * = outstanding reviewer

Journals: IJCV, TMLR, Pattern Recognition, CVIU.

Session Chair: WACV 2021, ICPR 2020.

SELECTED INTERNS AND THESISISTS SUPERVISED

Yura Choi (2025) - incoming Ph.D. student at Imperial College London.

Ye-Bin Moon (2025) - intern from POSTECH.

Xin Wen (2024-2025) - intern from the University of Hong Kong.

Aysim Toker (2024-2025) - intern from Technical University of Munich.

Tatiana Gaintseva (2024-2025) - intern from Queen Mary University.

Changrui Chen (2024) - intern from University of Warwick – > research scientist at Huawei.

Bingchen Zhao (2024-2025) - intern from University of Edinburgh.

Prannay Kaul (2024) - intern from University of Oxford – > research scientist at Amazon.

Yongshuo Zong (2024) - intern from University of Edninburgh, paper together in ICML workshop.

Konstantinos Alexandridis (2023) - intern from King's College – > research scientist at Huawei.

Roy Miles (2023) - intern from Imperial College, paper together at CVPR – > research scientist at Huawei.

Chengcheng Ma (2023) - intern from Chinese Academy of Sciences, paper at AAAI – > research scientist at Kunlun.

Yunqi Miao (2023) - intern at Huawei from University of Warwick – > research scientist at Huawei.

Jenny Seidenschwarz - masters and Ph.D. student at TUM (2020-2023), papers together at ICML, tPAMI, CVPR.

Franziska Gerken - Ph.D. student at TUM (2020-2023), submission together at eLife.

Volodymyr Fomenko - masters at TUM (2021-2022), paper together at NeurIPS – > Technical Staff at OpenAI.

Peter Kocsis - masters at TUM (2021-2022), paper together at NeurIPS – > Ph.D. student at TUM.

Laurin Wagner - masters at TUM (2020-2021), paper together at tPAMI – > ML Research Engineer at myReha.

Peter Sukenik - masters at TUM (2021), paper together at NeurIPS – > Ph.D. student at IST Austriat.

Feliks Hibraj (2020-2021) - intern at TUM – > software engineer at Snap Inc.

TEACHING EXPERIENCE

Deep Learning with Pytorch at Datacamp (2019): instructor. Developed during my Ph.D., over 28K students attended the course, before it got retired in December 2023.

Introduction to Deep Learning at TUM (2022): co-instructor. Gave half of the lectures, and was in charge of the exam. Around 1000 students attended the course.

Advanced Computer Vision at TUM (2021 and 2022): co-instructor. Gave several lectures, lead the office hours, and was in charge of the exam. 30 students attended the course.

Computer Vision III: Detection, Segmentation, and Tracking at TUM (2022): co-instructor. Gave several lectures, lead the office hours, and was in charge of the exam. Around 150 students attended the course.

Introduction to Machine Learning at Aralytics: instructor. Designed and gave a course for the company's internal training. 10 employees attended the course.